

EE516 : Homework 1

Gaurav Kalra

(Student ID: 2016 45 93)

gvkalra@kaist.ac.kr

Problem 1: Device Driver

Make your device driver with a Stack and its application program to prove the correctness.

FOLDERS

▼ driver

main.c

Makefile

stack.c

stack.h

utils.h

206 → module_init(stack_dev_init);

207 module_exit(stack_dev_exit); ←

208

209 MODULE_AUTHOR("Gaurav Kalra");

210 MODULE_DESCRIPTION("HW01 device driver providing a stack");

211 MODULE_LICENSE("GPL");

144 static int __init

145 stack_dev_init(void)

146 {

147 int ret;

148

149 dbg("");

150

151 /* Register char device number */

152 ret = alloc_chrdev_region(&dev_first, CHARDEV_NR_MINOR, CHARDEV_NR_DEVICES, CHARDEV_NAME);

153 if (ret < 0) {

154 err("Failed allocating device number");

155 goto error;

156 }

157

158 /* Allocate cdev structure */

159 cdev = cdev_alloc();

160 if (cdev == NULL) {

161 err("Failed allocating character device");

162 goto error;

163 }

164

165 cdev->owner = THIS_MODULE;

166 cdev->ops = &fops;

167

168 /* Add char device to the system */

169 ret = cdev_add(cdev, dev_first, CHARDEV_NR_DEVICES);

170 if (ret < 0) {

171 err("Failed adding character device to the system");

172 goto error;

173 }

174

175 info("[STACK_DEVICE] allocated Major(%d) and Minor(%d)", MAJOR(dev_first), MINOR(dev_first));

176

177 /* Create /dev/stack_device node

178 * Ref: https://github.com/euspectre/keedr/blob/master/sources/examples/sample_target/cfake.c

179 */

180 class = class_create(THIS_MODULE, CHARDEV_NAME);

181 if (class == NULL) {

182 err("Failed creating device class");

183 goto error;

184 }

185 device = device_create(class, NULL, dev_first, NULL, "%s", CHARDEV_NAME);

186 if (device == NULL) {

187 err("Failed creating /dev node");

188 goto error;

189 }

190

191 return 0;

192

193 error:

194 /* we can't call stack_dev_exit()

195 | because it's segment is controlled by __exit modifier

196 */

197 _stack_dev_exit();

198 return -1;

199 }

200 static void __exit

201 stack_dev_exit(void)

202 {

203 _stack_dev_exit();

204 }

Figure 1: Kernel Module initialization / de-initialization in main.c

```

112 static void
113 _stack_dev_exit(void)
114 {
115     dbg("");
116
117     /* Remove /dev/stack_device node */
118     if (device != NULL) {
119         device_destroy(class, dev_first);
120         device = NULL;
121     }
122
123     /* Destroy device class */
124     if (class != NULL) {
125         class_destroy(class);
126         class = NULL;
127     }
128
129     /* Remove cdev */
130     if (cdev != NULL) {
131         cdev_del(cdev);
132         cdev = NULL;
133     }
134
135     /* Un-register char device number */
136     if (dev_first != 0) {
137         unregister_chrdev_region(dev_first, CHARDEV_NR_DEVICES);
138         dev_first = 0;
139     }
140
141     info("[STACK_DEVICE] released");
142 }

```

Figure 1.1: Kernel Module initialization / de-initialization in main.c

A separate function `_stack_dev_exit()` is reused in `stack_dev_init()` and `stack_dev_exit()` because of different lifetimes provided by `__init` and `__exit` modifiers.

Major & Minor numbers are dynamically allocated by the Kernel using the new way of device allocation.

References:

<http://stackoverflow.com/questions/8563978/what-is-kernel-section-mismatch>
https://github.com/euspectre/kedr/blob/master/sources/examples/sample_target/cfake.c

```

98 static struct file_operations fops = {
99     .owner = THIS_MODULE,
100     .open = stack_dev_open, /* open() */
101     .read = stack_dev_read, /* read() */
102     .write = stack_dev_write, /* write() */
103     .release = stack_dev_release, /* close() */
104     /*
105      * References:
106      * [1] https://github.com/torvalds/linux/blob/master/fs/ext4/file.c#L698
107      * [2] https://github.com/torvalds/linux/blob/master/fs/ext4/ioctl.c#L436
108      */
109     .unlocked_ioctl = stack_dev_ioctl, /* unlocked_ioctl() */
110 };

```

```

26 static int
27 stack_dev_open(struct inode *inode, struct file *file)
28 {
29     dbg("");
30     return 0;
31 }

```

```

81 static int
82 stack_dev_release(struct inode *inode, struct file *file)
83 {
84     dbg("");
85     return 0;
86 }

```

```

88 static long
89 stack_dev_ioctl(struct file *file, unsigned int cmd, unsigned long arg)
90 {
91     dbg("");
92     st_clean(); 1
93     return 0;
94 }

```

```

33 static ssize_t
34 stack_dev_read(struct file *file, char __user *buffer, size_t length, loff_t *offset)
35 {
36     int item, ret;
37     dbg("");
38     /* EOF, return 0 */ 2
39     if (st_is_empty())
40         return 0;
41     /* pop */ 3
42     ret = st_pop(&item);
43     if (ret < 0)
44         return ret;
45     /* copy data into user space */ 4
46     if (copy_to_user(buffer, &item, sizeof(item)))
47         return -EFAULT; /* Bad address */
48     /* popped & copied to user space, return bytes */ 5
49     return sizeof(item);
50 }

```

```

57 static ssize_t
58 stack_dev_write(struct file *file, const char __user *buffer, size_t length, loff_t *offset)
59 {
60     int item, ret;
61     dbg("");
62     /* No more space */ 6
63     if (st_is_full())
64         return -ENOMEM;
65     /* copy data from user space */ 7
66     if (copy_from_user(&item, buffer, sizeof(item)))
67         return -EFAULT; /* Bad address */
68     /* push */ 8
69     ret = st_push(item);
70     if (ret < 0)
71         return ret;
72     /* pushed & copied from user space, return bytes */ 9
73     return sizeof(item);
74 }

```

Figure 2: VFS operations implemented by Stack device in main.c

```

1  #pragma once
2
3  /* True (1) if no more items can be popped and there is no top item
4  */
5  int st_is_empty(void);
6
7  /* True (1) if no more items can be pushed
8  */
9  int st_is_full(void);
10
11 /* Adds an item onto the stack
12  * Returns < 0 on error
13  */
14 int st_push(int item);
15
16 /* Removes the most-recently-pushed item from the stack (as out argument)
17  * Returns < 0 on error
18  */
19 int st_pop(int *item);
20
21 /* Cleans up stack
22  */
23 void st_clean(void);

```

Figure 3: APIs exposed by Stack abstraction in stack.h

```

6  #define STACK_SIZE 256
7
8  struct {
9      int stack[STACK_SIZE];
10     int top;
11 } st = {
12     .top = -1, /* top is set to -1 to indicate empty stack */
13 };

```

```

15 int st_is_empty(void)
16 {
17     dbg("");
18     if (st.top == -1)
19         return 1;
20     return 0;
21 }

```

```

33 int st_push(int item)
34 {
35     dbg("");
36     /* add item */
37     st.top++;
38     st.stack[st.top] = item;
39     return 0;
40 }

```

```

43 int st_pop(int *item)
44 {
45     /* invalid arguments */
46     if (item == NULL)
47         return -EINVAL;
48     dbg("");
49     /* remove item */
50     *item = st.stack[st.top];
51     st.top--;
52     return 0;
53 }

```

```

57 void st_clean(void)
58 {
59     dbg("");
60     /* setting top to -1 means stack is cleaned up
61     * there is no way to access st
62     */
63     st.top = -1;
64 }

```

```

24 int st_is_full(void)
25 {
26     dbg("");
27     if (st.top >= STACK_SIZE - 1)
28         return 1;
29     return 0;
30 }

```

Figure 4: Implementation of Stack abstraction in stack.c

```

1  MODNAME := stack_device
2  obj-m := ${MODNAME}.o
3  ${MODNAME}-objs := main.o stack.o
4
5  KDIR := /lib/modules/$(shell uname -r)/build
6
7  all:
8  | $(MAKE) -C $(KDIR) M=$(PWD) modules
9
10 clean:
11 | $(MAKE) -C $(KDIR) M=$(PWD) clean

```

Figure 5: Makefile for Kernel Module

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01/driver (master) $ make
make -C /lib/modules/4.4.0-38-generic/build M=/home/gvkalra/Desktop/EE516/HW01/driver modules
make[1]: Entering directory '/usr/src/linux-headers-4.4.0-38-generic'
CC [M] /home/gvkalra/Desktop/EE516/HW01/driver/main.o
CC [M] /home/gvkalra/Desktop/EE516/HW01/driver/stack.o
LD [M] /home/gvkalra/Desktop/EE516/HW01/driver/stack_device.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/gvkalra/Desktop/EE516/HW01/driver/stack_device.mod.o
LD [M] /home/gvkalra/Desktop/EE516/HW01/driver/stack_device.ko
make[1]: Leaving directory '/usr/src/linux-headers-4.4.0-38-generic'

```

Figure 6: Generating Kernel Module (*stack_device.ko*)

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01/driver (master) $ sudo insmod stack_device.ko
gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01/driver (master) $ dmesg
[17728.856884] <stack_dev_init:149>
[17728.856889] <stack_dev_init:174> [STACK_DEVICE] allocated Major(244) and Minor(0)

```

Figure 7: Inserting Kernel Module using *insmod*

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01/driver (master) $ lsmod | grep stack_device
stack_device      16384  0
gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01/driver (master) $ cat /proc/devices | grep stack_device
244 stack_device
gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01/driver (master) $ ls -al /dev/ | grep stack_device
crw----- 1 root root 244, 0 Sep 22 22:39 stack_device

```

Figure 8: Verifying module insertion using *lsmod*, */proc/devices* and creation of */dev/stack_device* character device. Note that Major(244) and Minor(0) have been assigned by the kernel at runtime.

```

7  #define STACK_SIZE 256
8
9  int main(int argc, const char *argv[])
10 {
11     int item = 0, fd;
12     ssize_t bytes_written;
13
14     /* open driver node */
15     fd = open("/dev/stack_device", O_WRONLY);
16     if (fd < 0) {
17         fprintf(stderr, "open() failed err: [%s]\n", strerror(errno));
18         return -1;
19     }
20
21     /* start pushing numbers to the driver */
22     while (item < STACK_SIZE) {
23         bytes_written = write(fd, &item, sizeof(item));
24
25         /* error */
26         if (bytes_written < 0) {
27             fprintf(stderr, "write() failed err: [%s]\n", strerror(errno));
28         }
29         /* partial bytes written */
30         else if (bytes_written != sizeof(item)) {
31             fprintf(stderr, "[NOT WRITTEN] %d\n", item);
32         }
33         /* success */
34         else {
35             fprintf(stdout, "[WRITTEN] %d\n", item);
36             item++; /* push next item */
37         }
38     }
39
40     /* cleanup */
41     close(fd);
42     return 0;
43 }

```

Figure 9: *app1* writing 256 items (0 – 255) on *stack_device*

```

7  int main(int argc, const char *argv[])
8  {
9      int item, fd;
10     ssize_t bytes_read;
11
12     /* open driver node */
13     fd = open("/dev/stack_device", O_RDONLY);
14     if (fd < 0) {
15         fprintf(stderr, "open() failed err: [%s]\n", strerror(errno));
16         return -1;
17     }
18
19     /* start popping numbers from the driver */
20     while (1) {
21         bytes_read = read(fd, &item, sizeof(item));
22
23         /* error */
24         if (bytes_read < 0) {
25             fprintf(stderr, "read() failed err: [%s]\n", strerror(errno));
26         }
27         /* EOF */
28         else if (bytes_read == 0) {
29             break;
30         }
31         /* incomplete data */
32         else if (bytes_read != sizeof(item)) {
33             fprintf(stderr, "[NOT READ] %d\n", item);
34         }
35         /* success */
36         else {
37             fprintf(stdout, "[READ] %d\n", item);
38         }
39     }
40
41     /* cleanup */
42     close(fd);
43     return 0;
44 }

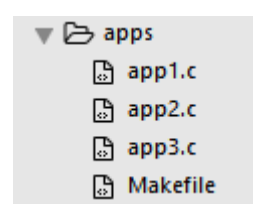
```

Figure 10: *app2* reading items until EOF from *stack_device*

```

1 CC=gcc
2 CFLAGS=-Wall -Werror
3 BIN=app1 app2 app3
4
5 all: $(BIN)
6
7 app1:
8     $(CC) -o $@ $.c $(CFLAGS)
9
10 app2:
11     $(CC) -o $@ $.c $(CFLAGS)
12
13 app3:
14     $(CC) -o $@ $.c $(CFLAGS)
15
16 clean:
17     rm -f $(BIN)

```



```

gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01 (master) $ make
gcc -o app1 app1.c -Wall -Werror
gcc -o app2 app2.c -Wall -Werror
gcc -o app3 app3.c -Wall -Werror

```

Figure 11: Makefile & folder structure of applications

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01 (master) $ sudo ./app1 | head -n 10
[WRITTEN] 0
[WRITTEN] 1
[WRITTEN] 2
[WRITTEN] 3
[WRITTEN] 4
[WRITTEN] 5
[WRITTEN] 6
[WRITTEN] 7
[WRITTEN] 8
[WRITTEN] 9

```

Figure 12: Output of ./app1 (write) stripped to 10 lines

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01 (master) $ sudo ./app2 | head -n 10
[READ] 255
[READ] 254
[READ] 253
[READ] 252
[READ] 251
[READ] 250
[READ] 249
[READ] 248
[READ] 247
[READ] 246

```

Figure 13: Output of ./app2 (read) stripped to 10 lines

The read data sequence is opposite of write data sequence because the *read()* & *write()* system calls on */dev/stack_device* are implemented as a stack data structure (first-in last-out, last-in first-out) in our driver.

Problem 2: Clean Function

Make your own “clean function” that makes the stack empty.

```
8  int main(int argc, const char *argv[])
9  {
10     int fd, ret;
11
12     /* open driver node */
13     fd = open("/dev/stack_device", O_RDONLY);
14     if (fd < 0) {
15         fprintf(stderr, "open() failed err: [%s]\n", strerror(errno));
16         return -1;
17     }
18
19     /* call ioctl() on driver */
20     ret = ioctl(fd, 0);
21     if (ret < 0) {
22         fprintf(stderr, "ioctl() failed err: [%s]\n", strerror(errno));
23         close(fd);
24         return -1;
25     }
26
27     /* cleanup */
28     close(fd);
29     return 0;
30 }
```

Figure 14: *app3* cleans *stack_device* by invoking *ioctl()*

The structure *file_operations* provides an extension point for custom commands using *ioctl* (I/O control). I have implemented the operation of “stack clean” in Linux using *unlocked_ioctl()* [**Refer Figure 2**].

It can be invoked by an application using *ioctl()* system call.

Reference:

<https://github.com/torvalds/linux/blob/master/include/linux/fs.h#L1679>

```
gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01 (master) $ sudo ./app1 | head -n 10
[WRITTEN] 0
[WRITTEN] 1
[WRITTEN] 2
[WRITTEN] 3
[WRITTEN] 4
[WRITTEN] 5
[WRITTEN] 6
[WRITTEN] 7
[WRITTEN] 8
[WRITTEN] 9
gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01 (master) $ sudo ./app3
gvkalra@gvkalra-desktop ~/Desktop/EE516/HW01 (master) $ sudo ./app2
```

Figure 15: *app1* writes 256 items (0 – 255) on *stack_device*, *app3* cleans *stack_device*, *app2* has no remaining data to read.