

EE516 : Project 1

Gaurav Kalra

(Student ID: 2016 45 93)

gvkalra@kaist.ac.kr

Task 1: Tree Script

Write a bash shell script that shows tree structure of the directories and files included in your own home directory.

```
1  #!/bin/bash
2
3  # set bash options
4  # Reference: https://www.gnu.org/software/bash/manual/html_node/The-Shopt-Builtin.html
5  #
6  # dotglob:
7  #   If set, Bash includes filenames beginning with a '.' in the results of filename expansion.
8  #   (for hidden files)
9  #
10 # nullglob:
11 #   If set, Bash allows filename patterns which match no files to expand to a null string,
12 #   rather than themselves. (for empty directories)
13 shopt -s dotglob nullglob 1
14
15 # For tracking the depth of file
16 DEPTH=0
17
18 # Function: print_tree
19 # Purpose: To print tree structure of a folder in filesystem
20 # Arguments
21 #   ($1): Folder to be used as root node of tree
22 print_tree() { 2
23     cd "$1"
24
25     # for each file
26     for FILE in * 3
27     do
28         # print "|" between each depth
29         ITER=0
30         while [ $ITER != $DEPTH ]
31         do
32             echo -n "|"
33             ITER=$((ITER + 1))
34         done
35
36         # print file
37         echo -n "|____" 4
38         echo $FILE
39
40         # recurse if file is directory
41         if [ -d "$FILE" ]; then
42             DEPTH=$((DEPTH + 1))
43             print_tree "$FILE" 5
44             cd ..
45         fi
46     done
47
48     # move to previous depth & process more files 6
49     DEPTH=$((DEPTH - 1));
50 }
51
52 # For aesthetics
53 echo "|____"
54
55 # Print tree structure rooted at home
56 print_tree ~ 7
57
58 # unset bash options
59 shopt -u dotglob nullglob
60
61 # Return success (0) for $?
62 exit 0
```

Figure 1: Source code of Tree Script with Key Points highlighted

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ ls -al | grep tree
-rwxrwxr-x 1 gvkalra gvkalra 1218 Sep 15 22:08 tree.sh
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ ./tree.sh | head -n 20
|
|___ .bash_history
|___ .bash_logout
|___ .bash_profile
|___ .bashrc
|___ .cache
|___ compizconfig-1
|   |___ animation.pb
|   |___ commands.pb
|   |___ compiztoolbox.pb
|   |___ composite.pb
|   |___ copytex.pb
|   |___ core.pb
|   |___ decor.pb
|   |___ expo.pb
|   |___ ezoom.pb
|   |___ fade.pb
|   |___ gnomecompat.pb
|   |___ grid.pb
|   |___ imgpng.pb

```

Figure 2: Sample output of executing Tree Script stripped to 20 lines

The core of *tree.sh* is implemented as a recursive function *print_tree()*, which keeps track of the depth of a sub-directory using *DEPTH* variable.

Bash options *dotglob* and *nullglob* are necessary to be set for enabling *for FILE in ** loop iterate on hidden files & reporting contents of an empty directory as null string respectively.

It is also important to double quote *\$1* and *\$FILE* since the filenames may contain spaces, in which case *print_tree()* will treat one filename as multiple arguments.

```

207 static void
208 file_read_sequential(void)
209 {
210     int i, fd;
211     char filename[128], *buf;
212     ssize_t bytes_read;
213
214     1 if ((FILESIZE % req_size) != 0) {
215         err("FILESIZE(%d) and req_size(%d) are not aligned",
216             FILESIZE, req_size);
217         exit(1);
218     }
219
220     buf = memalign((size_t)req_size, (size_t)req_size);
221     if (buf == NULL) {
222         err("Failed to allocate buffer");
223         exit(1);
224     }
225
226     for (i = 0; i < NUMFILES; i++) {
227         snprintf(filename, 128, "%s/file-%d", dirname, i);
228         fd = open(filename, O_RDONLY);
229         if (fd == -1) {
230             err("open() failed: [%s]", strerror(errno));
231             free(buf);
232             exit(1);
233         }
234         info("File Opened Sequential Read ..");
235
236         2 do {
237             /* read chunks of req_size bytes */
238             bytes_read = read(fd, buf, req_size);
239             if (bytes_read == -1) {
240                 err("read() failed: [%s]", strerror(errno));
241                 free(buf);
242                 close(fd);
243                 exit(1);
244             }
245             /* read until EOF is reached */
246             } while (bytes_read != 0);
247
248         close(fd);
249     }
250     free(buf);
251 }

```

Figure 3: Source code of Sequential Read with Key Points highlighted

The check for *FILESIZE % req_size* to be zero is added as a sanity measure for invalid inputs.

For sequential *read()*, I try to read *req_size* bytes of data in a while loop until EOF is reached.

```

253 static void
254 file_write_random(void)
255 {
256     int i, offset, fd;
257     char filename[128], *buf;
258     ssize_t bytes_written;
259
260     1 if ((FILESIZE % req_size) != 0) {
261         err("FILESIZE(%d) and req_size(%d) are not aligned",
262             FILESIZE, req_size);
263         exit(1);
264     }
265
266     buf = memalign((size_t)req_size, (size_t)req_size);
267     if (buf == NULL) {
268         err("Failed to allocate buffer");
269         exit(1);
270     }
271
272     for (i = 0; i < NUMFILES; i++) {
273         snprintf(filename, 128, "%s/file-%d", dirname, i);
274         fd = open(filename, O_DIRECT | O_WRONLY | O_EXCL,
275             S_IWUSR | S_IRUSR | S_IRGRP | S_IWGRP | S_IROTH | S_IWOTH);
276         if (fd == -1) {
277             err("open() failed: [%s]", strerror(errno));
278             free(buf);
279             exit(1);
280         }
281         info("File Opened Random Write ..");
282
283         /* randomly get an offset for lseek() */
284         offset = get_next_rand_number((FILESIZE / req_size));
285         while (offset >= 0) {
286             /* seek from beginning of file */
287             lseek(fd, (offset * req_size), SEEK_SET);
288
289             /* at every random offset, flush req_size bytes of data
290              *
291              * As noted in file_write_sequential(), it is assumed that write()
292              * will either flush whole req_size bytes of data or none at all.
293              * The case where physical medium is full & there is not enough space
294              * is not handled
295              */
296             bytes_written = write(fd, buf, req_size);
297             if (bytes_written == -1) {
298                 err("write() failed: [%s]", strerror(errno));
299                 free(buf);
300                 close(fd);
301                 exit(1);
302             }
303
304             /* generate another random offset */
305             offset = get_next_rand_number(-1);
306         }
307         close(fd);
308     }
309     free(buf);
310 }

```

Figure 4: Source code of Random Write with Key Points highlighted

For random *write()*, I generate a unique random number, *lseek()* to the random offset & try to write *req_size* bytes of data until random number pool is exhausted (size of file becomes *FILESIZE*).

```

312 static void
313 file_read_random(void)
314 {
315     int i, fd, offset;
316     char filename[128], *buf;
317     ssize_t total_bytes, bytes_read;
318
319     1 if ((FILESIZE % req_size) != 0) {
320         err("FILESIZE(%d) and req_size(%d) are not aligned",
321             FILESIZE, req_size);
322         exit(1);
323     }
324
325     buf = memalign((size_t)req_size, (size_t)req_size);
326     if (buf == NULL) {
327         err("Failed to allocate buffer");
328         exit(1);
329     }
330
331     for (i = 0; i < NUMFILES; i++) {
332         snprintf(filename, 128, "%s/file-%d", dirname, i);
333         fd = open(filename, O_RDONLY);
334         if (fd == -1) {
335             err("open() failed: [%s]", strerror(errno));
336             free(buf);
337             exit(1);
338         }
339         info("File Opened Random Read ..");
340
341         /* randomly get an offset for lseek() */
342         offset = get_next_rand_number((FILESIZE / req_size));
343         while (offset >= 0) {
344             /* seek from beginning of file */
345             lseek(fd, (offset * req_size), SEEK_SET);
346
347             total_bytes = 0;
348             do {
349                 bytes_read = read(fd, buf, req_size);
350                 if (bytes_read == -1) {
351                     err("read() failed: [%s]", strerror(errno));
352                     free(buf);
353                     close(fd);
354                     exit(1);
355                 }
356
357                 /* at every random offset, read req_size bytes of data
358                  *
359                  * In case read() is not able to read whole req_size bytes,
360                  * we retry to read until total_bytes for the current random offset
361                  * reach req_size. This is best effort basis to sequentially
362                  * read the whole file.
363                  */
364                 total_bytes += bytes_read;
365             } while (total_bytes != req_size);
366
367             /* generate another random offset */
368             offset = get_next_rand_number(-1);
369         }
370         close(fd);
371     }
372     free(buf);
373 }

```

Figure 5: Source code of Random Read with Key Points highlighted

```

57  /* Generates a unique (non-repeating) random number in O(1)
58  *
59  * pool_size is how many unique numbers are required to be generated.
60  *
61  * e.g.
62  * The function get_next_rand_number(512) will initialize a pool of 512 random numbers (0 to 511)
63  * and return a random number from it.
64  * Subsequent calls to get_next_rand_number(-1) will return a unique random number from the pool
65  * A call to get_next_rand_number(-1) will return -1 when the pool is exhausted
66  *
67  * Ref: http://stackoverflow.com/questions/196017/unique-non-repeating-random-numbers-in-o1
68  */
69  static int
70  get_next_rand_number(int pool_size)
71  {
72      int i, num, temp;
73      static int num_left = 0;
74      static int *num_pool = NULL;
75
76      /* Initialize a new pool */
77      if (pool_size >= 0) {
78          /* A pool already exists. Free it first */
79          if (num_pool != NULL) {
80              free(num_pool);
81              num_pool = NULL;
82          }
83
84          /* Allocate memory */
85          num_pool = malloc(sizeof(int) * pool_size);
86          if (num_pool == NULL) {
87              err("malloc() failed: [%s]", strerror(errno));
88              exit(1);
89          }
90
91          /* Initial values */
92          for (i = 0; i < pool_size; i++)
93              num_pool[i] = i;
94          num_left = pool_size;
95      }
96
97      /* pool exhausted, return -1 */
98      if (num_left == 0)
99          return -1;
100
101      /* select a random number between 0 and num_left */
102      num = rand() % num_left;
103
104      /* replace num_pool[num] with num_pool[num_left - 1] */
105      temp = num_pool[num];
106      num_pool[num] = num_pool[num_left - 1];
107
108      /* decrement items left in pool */
109      num_left--;
110
111      /* if pool is exhausted, free memory */
112      if (num_left == 0) {
113          free(num_pool);
114          num_pool = NULL;
115      }
116
117      /* return previous num_pool[num] (saved in temp) */
118      return temp;
119  }

```

Figure 6: Random Number Generator with Key Points highlighted

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ ls -al | grep fsbench
-rw-rw-r-- 1 gvkalra gvkalra 10018 Sep 17 20:41 fsbench.c
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ gcc -o fsbench fsbench.c -Wall -Werror
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ ls -al | grep fsbench
-rwxrwxr-x 1 gvkalra gvkalra 18688 Sep 18 19:51 fsbench
-rw-rw-r-- 1 gvkalra gvkalra 10018 Sep 17 20:41 fsbench.c
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ ./fsbench ./ 1024
File Created ..
File Opened Sequential Write ..
File Opened Sequential Read ..
File Opened Random Write ..
File Opened Random Read ..
===== File System Benchmark Execution Result (Time usec) =====
File Create      :      123
Sequential Write :    41989
Sequential Read  :    11032
Random Write     :    62068
Random Read     :    18684
File Delete      :     108
Total           :   134004
=====
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ ./fsbench ./ 2048
File Created ..
File Opened Sequential Write ..
File Opened Sequential Read ..
File Opened Random Write ..
File Opened Random Read ..
===== File System Benchmark Execution Result (Time usec) =====
File Create      :     132
Sequential Write :   38747
Sequential Read  :   10928
Random Write     :   31654
Random Read     :   14430
File Delete      :     136
Total           :   96027
=====
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ ./fsbench ./ 2044
File Created ..
File Opened Sequential Write ..
<file_write_sequential:195> write() failed: [Invalid argument]

```

Figure 7: Output of compiling & executing File System benchmark

Please notice that `./fsbench ./ 2044` fails because current file offset, buffer, bytes to be written are not aligned in `file_write_sequential()`. Subsequent executions of `fsbench` will as well fail until “`rm -f file-*`” is executed in the current directory. This is a known & easy to fix bug. However, it is outside the purview of current assignment.

Notice that there are no warnings (source is compiled with `-Wall -Werror` flags).


```

gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ gcc -g -o fsbench fsbench.c -Wall -Werror
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ valgrind ./fsbench ./ 1024
==9071== Memcheck, a memory error detector
==9071== Copyright (C) 2002-2015, and GNU GPL'd, by Julian Seward et al.
==9071== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==9071== Command: ./fsbench ./ 1024
==9071==
File Created ..
File Opened Sequential Write ..
==9071== Syscall param write(buf) points to uninitialised byte(s)
==9071==    at 0x4F30A10: __write_nocancel (syscall-template.S:84)
==9071==    by 0x4010BD: file_write_sequential (fsbench.c:193)
==9071==    by 0x401A48: main (fsbench.c:391)
==9071== Address 0x5203800 is 0 bytes inside a block of size 1,024 alloc'd
==9071==    at 0x4C2FFC6: memalign (vg_replace_malloc.c:858)
==9071==    by 0x400F75: file_write_sequential (fsbench.c:165)
==9071==    by 0x401A48: main (fsbench.c:391)
==9071==
File Opened Sequential Read ..
File Opened Random Write ..
==9071== Syscall param write(buf) points to uninitialised byte(s)
==9071==    at 0x4F30A10: __write_nocancel (syscall-template.S:84)
==9071==    by 0x40160F: file_write_random (fsbench.c:296)
==9071==    by 0x401A6E: main (fsbench.c:400)
==9071== Address 0x5204800 is 0 bytes inside a block of size 1,024 alloc'd
==9071==    at 0x4C2FFC6: memalign (vg_replace_malloc.c:858)
==9071==    by 0x401491: file_write_random (fsbench.c:266)
==9071==    by 0x401A6E: main (fsbench.c:400)
==9071==
File Opened Random Read ..
===== File System Benchmark Execution Result (Time usec) =====
File Create      :      12098
Sequential Write :      56784
Sequential Read  :      18434
Random Write     :      82319
Random Read     :      26845
File Delete      :       1118
Total           :     197598
=====
==9071==
==9071== HEAP SUMMARY:
==9071==    in use at exit: 0 bytes in 0 blocks
==9071==   total heap usage: 7 allocs, 7 frees, 13,312 bytes allocated
==9071==
==9071== All heap blocks were freed -- no leaks are possible
==9071==
==9071== For counts of detected and suppressed errors, rerun with: -v
==9071== Use --track-origins=yes to see where uninitialised values come from
==9071== ERROR SUMMARY: 2048 errors from 2 contexts (suppressed: 0 from 0)

```

Figure 8: Output of compiling & executing File System benchmark with valgrind for checking heap memory leaks

Please notice “HEAP SUMMARY” which says “All heap blocks were freed – no leaks are possible”

Task 3: Makefile

Write a Makefile for the source code of Task 2

```
1 CC=gcc
2 CFLAGS=-Wall -Werror
3 BIN=fsbench
4
5 all: $(BIN)
6
7 $(BIN):
8     $(CC) -o $@ $@.c $(CFLAGS)
9
10 clean:
11     rm -f $(BIN)
```

Figure 9: Makefile of Task 2 with Key Points highlighted

```
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ make
gcc -o fsbench fsbench.c -Wall -Werror
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ make clean
rm -f fsbench
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ make all
gcc -o fsbench fsbench.c -Wall -Werror
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR01 (master) $ make clean
rm -f fsbench
```

Figure 10: Sample output of Makefile targets