

EE516 : Project 2

Gaurav Kalra

(Student ID: 2016 45 93)

gvkalra@kaist.ac.kr

Task 1: Analyze the given skeleton program

Find the functions that you do not know and explain the operations of those functions

```
66 static int __init init_procmon (void)
67 {
68
69     struct proc_dir_entry *procmon_proc;
70
71     1 procmon_proc = proc_create(FILE_NAME, 644, NULL, &fops);
72     if (!procmon_proc) {
73         printk(KERN_ERR "Cannot create procmon proc entry \n");
74         return -1;
75     }
76     printk(KERN_INFO "init procmon\n");
77     return 0;
78 }
79
80
81 static void __exit exit_procmon(void)
82 {
83     2 remove_proc_entry(FILE_NAME, NULL);
84     printk(KERN_INFO "exit procmon\n");
85 }
86
87
88
89 module_init(init_procmon);
90 module_exit(exit_procmon);
91
92 MODULE_LICENSE("GPL");
93 MODULE_AUTHOR("Dong-Jae Shin");
94 MODULE_DESCRIPTION("EE516 Project2 Process Monitoring Module");
95
```

proc_create(): creates a file entry in /proc directory
parameters:

- 1. **name** – filename which will show up in proc directory
- 2. **mode** – file creation mode (w,r,x for user,group,others)
- 3. **parent directory** – pointer to parent directory in /proc where this file will be created (we used NULL to indicate /proc/ directory)
- 4. **file operations** – pointer to file operations structure

remove_proc_entry(): remove the file entry from /proc directory
parameters:

- 1. **name** – filename as in /proc
- 2. **parent** – parent dir pointer (same NULL here)

Figure 1: Creating / Removing a file from procfs

```
56 struct file_operations fops = {
57     .owner = THIS_MODULE,
58     .open = procmon_proc_open,
59     1 .read = seq_read,
60     .llseek = seq_lseek,
61     2 .write = procmon_proc_write,
62     .release = single_release,
63 };
```

seq_read(): read method for sequential files
seq_lseek(): llseek method for sequential files
single_release(): free the structures associated with sequential files

The parameters for these APIs are compatible with respective VFS file operations. And thus can be used as drop-in replacement so that we do not need to write our own implementation.

Figure 2: Sequence File APIs for VFS file operations

```
28 static int procmon_proc_show(struct seq_file *m, void *v)
29 {
30     1 seq_printf(m, "==== Contents ===== \n");
31     return 0;
32 }
33
34
35
36 static int procmon_proc_open(struct inode *inode, struct file *file)
37 {
38     2 return single_open(file, procmon_proc_show, NULL);
39 }
40
```

seq_printf(): It works like printk() for proc files. The first argument must be the sequence file.

single_open(): It creates and opens a new file descriptor for sequence files.

parameters:

- 1. **file** – The file structure from VFS
- 2. **show** – Call-back for printing data in proc filesystem
- 3. **data** – Call-back data (NULL in our case)

```
36 static ssize_t procmon_proc_write(struct file *file, const char __user *buf, size_t count, loff_t *ppos)
37 {
38
39     memset(mybuf, 0, sizeof(mybuf));
40
41     if (count > BUF_SIZE) {
42         count = BUF_SIZE;
43     }
44
45     if (copy_from_user(mybuf, buf, count)) {
46         return -EFAULT;
47     }
48
49     printk(KERN_INFO "proc write : %s\n", mybuf);
50     return (ssize_t)count;
51 }
52
```

Figure 3: Using Sequence Files for displaying contents of procfs

Task 2: Traverse Process - tasklist

Print every task's PID and Process Name in your proc file system

```
57  /* file operations */
58  static struct file_operations fops = {
59      .owner = THIS_MODULE,
60      .open = pl_open,
61
62      /* read method for sequential files */
63      .read = seq_read,
64
65      /* llseek method for sequential files */
66      .llseek = seq_lseek,
67
68      /* free the structures associated with sequential file */
69      .release = single_release,
70  };
71
72  static void
73  _pl_module_exit(void)
74  {
75      dbg("");
76
77      if (pl != NULL)
78          proc_remove(pl);
79  }
80
81  static int __init
82  pl_module_init(void)
83  {
84      dbg("");
85
86      /* create /proc/proc_list */
87      pl = proc_create(PROC_NAME, 0, NULL, &fops);
88      if (pl == NULL) {
89          err("Failed to create proc_list");
90          goto error;
91      }
92      return 0;
93
94  error:
95      _pl_module_exit();
96      return -1;
97  }
98
99  static void __exit
100 pl_module_exit(void)
101 {
102     _pl_module_exit();
103 }
104
105 module_init(pl_module_init);
106 module_exit(pl_module_exit);
107
108 MODULE_AUTHOR("Gaurav Kalra");
109 MODULE_DESCRIPTION("PR02 Traverse Process - tasklist");
110 MODULE_LICENSE("GPL");
```

Figure 4: Module Initialization / Clean-up (refer task01 for details)

```

12  /*
13  Organization of task information in kernel:
14
15      struct task_struct {
16          ...
17          pid_t pid;
18          ...
19          struct list_head tasks;
20          ...
21          char comm[TASK_COMM_LEN];
22      };
23
24      struct list_head {
25          struct list_head *next, *prev;
26      };
27  */
28
29  static int
30  pl_show(struct seq_file *m, void *v)
31  {
32      struct task_struct *tsk;
33      char name[TASK_COMM_LEN];
34
35      /* header */
36      seq_printf(m, "PID      ProcessName      \n");
37
38      /* Print pid & name of each process */
39      for_each_process(tsk) {
40
41          /* print information */
42          seq_printf(m, "%-10u%-20s\n",
43                  task_pid_nr(tsk),
44                  get_task_comm(name, tsk));
45      }
46
47      return 0;
48  }
49
50  static int
51  pl_open(struct inode *inode, struct file *file)
52  {
53      dbg("");
54      return single_open(file, pl_show, NULL);
55  }
56

```

Figure 5: Logic for printing PID & Process Name

```

1  MODNAME := proc_list
2  obj-m := ${MODNAME}.o
3  ${MODNAME}-objs := main.o
4
5  KDIR := /lib/modules/${shell uname -r}/build
6
7  all:
8      $(MAKE) -C $(KDIR) M=$(PWD) modules
9
10 clean:
11     $(MAKE) -C $(KDIR) M=$(PWD) clean

```

```

gvk@gvk:~/Desktop/EE516/PR02/task02 (master) $ make
make -C /lib/modules/4.4.0-38-generic/build M=/home/gvkalra/Desktop/EE516/PR02/task02 modules
make[1]: Entering directory '/usr/src/linux-headers-4.4.0-38-generic'
CC [M] /home/gvkalra/Desktop/EE516/PR02/task02/main.o
LD [M] /home/gvkalra/Desktop/EE516/PR02/task02/proc_list.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/gvkalra/Desktop/EE516/PR02/task02/proc_list.mod.o
LD [M] /home/gvkalra/Desktop/EE516/PR02/task02/proc_list.ko
make[1]: Leaving directory '/usr/src/linux-headers-4.4.0-38-generic'

```

Figure 6: Makefile (no warnings)

```

gvk@gvk:~/Desktop/EE516/PR02/task02 (master) $ cat /proc/proc_list | wc -l
290
gvk@gvk:~/Desktop/EE516/PR02/task02 (master) $ ps -ef | wc -l
290

```

Figure 7: Verification (proc_list & ps -ef both output 290 processes)

Task 3: per Process Memory Usage

Print every task's VIRT and RSS Memory Size

```
13  /*
14  Organization of task information in kernel:
15
16      struct task_struct {
17          ...
18          pid_t pid;
19          ...
20          struct list_head tasks;
21          ...
22          char comm[TASK_COMM_LEN];
23          ...
24          struct mm_struct *active_mm;
25      };
26
27      struct list_head {
28          struct list_head *next, *prev;
29      };
30
31      struct mm_struct {
32          ...
33          unsigned long total_vm;
34          ...
35          struct mm_rss_stat rss_stat;
36      };
37
38      struct mm_rss_stat {
39          atomic_long_t count[NR_MM_COUNTERS];
40      };
41
42      enum {
43          MM_FILEPAGES, //Resident file mapping pages (Type of File Mapped Page)
44          MM_ANONPAGES, //Resident anonymous pages (Type of Anonymous Page - Stack, Heap)
45          MM_SWAPENTS,  //Anonymous swap entries
46          MM_SHMEMPAGES, //Resident shared memory pages
47          NR_MM_COUNTERS
48      };
49
50      MM_FILEPAGES + MM_ANONPAGES = RSS Memory
51      total_vm = VIRT Memory
52  */
53
```

Figure 8: Organization of VIRT & RSS in task_struct

```

54 static int
55 pl_show(struct seq_file *m, void *v)
56 {
57     struct task_struct *tsk;
58     char name[TASK_COMM_LEN];
59     unsigned long long virt = 0;
60     long long rss = 0;
61
62     /* header */
63     seq_printf(m, "PID          ProcessName          "
64                "VIRT(KB)          RSS Mem(KB)          \n");
65
66     /* Print name, PID, VIRT & RSS of each process */
67     for_each_process(tsk) {
68         virt = rss = 0;
69
70         /* It is possible for active_mm to be NULL.
71          * In this case, we simply assume VIRT and RSS to be 0.
72          * Ref: https://www.kernel.org/doc/Documentation/vm/active\_mm.txt
73          */
74         if (tsk->active_mm != NULL) {
75             /* VIRT memory */
76             virt = tsk->active_mm->total_vm;
77             virt *= (PAGE_SIZE >> 10); /* convert pages to KB */
78
79             /* RSS Memory */
80             rss += atomic_long_read(&tsk->active_mm->rss_stat.count[MM_FILEPAGES]);
81             rss += atomic_long_read(&tsk->active_mm->rss_stat.count[MM_ANONPAGES]);
82             rss *= (PAGE_SIZE >> 10); /* convert pages to KB */
83         }
84
85         /* print information */
86         seq_printf(m, "%-10u%-20s%-20llu%-20llu\n",
87                    task_pid_nr(tsk),
88                    get_task_comm(name, tsk),
89                    virt, rss);
90     }
91
92     return 0;
93 }
94

```

Figure 9: Logic for calculating and printing VIRT, RSS

```

1  MODNAME := proc_memory
2  obj-m := ${MODNAME}.o
3  ${MODNAME}-objs := main.o
4
5  KDIR := /lib/modules/$(shell uname -r)/build
6
7  all:
8      $(MAKE) -C $(KDIR) M=$(PWD) modules
9
10 clean:
11     $(MAKE) -C $(KDIR) M=$(PWD) clean

```

Figure 10: Makefile (no warnings)

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/PR02/task03 (master) $ make
make -C /lib/modules/4.4.0-38-generic/build M=/home/gvkalra/Desktop/EE516/PR02/task03 modules
make[1]: Entering directory '/usr/src/linux-headers-4.4.0-38-generic'
CC [M] /home/gvkalra/Desktop/EE516/PR02/task03/main.o
LD [M] /home/gvkalra/Desktop/EE516/PR02/task03/proc_memory.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/gvkalra/Desktop/EE516/PR02/task03/proc_memory.mod.o
LD [M] /home/gvkalra/Desktop/EE516/PR02/task03/proc_memory.ko
make[1]: Leaving directory '/usr/src/linux-headers-4.4.0-38-generic'

```

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/PR02/task03 (master) $ cat /proc/proc_memory | grep sublime
16199  sublime_text 1096892 67768
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR02/task03 (master) $ top -b -n 1 | grep sublime
16199 gvkalra 26 0 1096892 67768 52368 S 0.0 0.8 0:00.44 sublime_text

```

Figure 11: Verification (sublime_text)

```

gvkalra@gvkalra-desktop ~/Desktop/EE516/PR02/task03 (master) $ top -b -n 1 | grep slack
2529 gvkalra 20 0 1686468 146920 85420 S 0.0 1.8 3:36.89 slack
2600 gvkalra 20 0 316428 30296 27184 S 0.0 0.4 0:00.01 slack
2669 gvkalra 20 0 1102600 104920 61124 S 0.0 1.3 0:08.41 slack
2736 gvkalra 20 0 1496056 346040 224816 S 0.0 4.3 1:57.91 slack
2780 gvkalra 20 0 1379216 265472 136304 S 0.0 3.3 0:45.19 slack
gvkalra@gvkalra-desktop ~/Desktop/EE516/PR02/task03 (master) $ cat /proc/proc_memory | grep slack
2529  slack 1686468 146920
2600  slack 316428 30296
2669  slack 1102600 104920
2736  slack 1496056 346040
2780  slack 1379216 265472

```

Figure 12: Verification (slack)

Task 4: per Process I/O Usage

Print every task's I/O usage

```
13  /*
14  Organization of task information in kernel:
15
16  struct task_struct {
17      ...
18      pid_t pid;
19      ...
20      struct list_head tasks;
21      ...
22      char comm[TASK_COMM_LEN];
23      ...
24      struct task_io_accounting ioac;
25  };
26
27  struct list_head {
28      struct list_head *next, *prev;
29  };
30
31  struct task_io_accounting {
32      #ifdef CONFIG_TASK_XACCT
33          u64 rchar; //bytes read
34          u64 wchar; //bytes written
35          u64 syscr; //# of read syscalls
36          u64 syscw; //# of write syscalls
37      #endif
38
39      #ifdef CONFIG_TASK_IO_ACCOUNTING
40          //The number of bytes which this task has caused to be read from storage.
41          u64 read_bytes;
42
43          //The number of bytes which this task has caused, or shall cause to be written to disk.
44          u64 write_bytes;
45
46          //A task can cause "negative" IO too. If this task truncates some
47          //dirty pagecache, some IO which another task has been accounted for
48          //(in its write_bytes) will not be happening. We _could_ just
49          //subtract that from the truncating task's write_bytes, but there is
50          //information loss in doing that.
51          u64 cancelled_write_bytes;
52      #endif
53  };
54  */
```

Figure 13: Organization of I/O data in task_struct


```

56 static int
57 pl_show(struct seq_file *m, void *v)
58 {
59     struct task_struct *tsk, *t;
60     char name[TASK_COMM_LEN];
61     struct task_io_accounting acct;
62
63     /* header */
64     seq_printf(m, "PID      ProcessName      "
65         "rchar(B)      wchar(B)      "
66         "syscr(#)      syscw(#)      "
67         "read_bytes(B)  write_bytes(B)  "
68         "cancelled_write_bytes(B)\n");
69
70     /* Print name, PID, I/O stats of each process */
71     for_each_process(tsk) {
72         acct = tsk->ioac; /* initialize accounting data */ 1
73
74         /* account each thread
75          * Ref: https://github.com/torvalds/linux/blob/master/fs/proc/base.c
76          * Function: do_io_accounting()
77          */
78         t = tsk;
79         task_io_accounting_add(&acct, &tsk->signal->ioac); 2
80         while_each_thread(tsk, t)
81             task_io_accounting_add(&acct, &t->ioac);
82
83         /* print information */
84         seq_printf(m, "%-10u%-20s%-20llu%-20llu%-20llu%-20llu%-20llu%-24llu\n",
85             task_pid_nr(tsk),
86             get_task_comm(name, tsk),
87             (unsigned long long)acct.rchar,
88             (unsigned long long)acct.wchar,
89             (unsigned long long)acct.syscr,
90             (unsigned long long)acct.syscw,
91             (unsigned long long)acct.read_bytes,
92             (unsigned long long)acct.write_bytes,
93             (unsigned long long)acct.cancelled_write_bytes);
94     }
95
96     return 0;
97 }

```

Figure 14: Logic for calculating and printing I/O information

```

1  MODNAME := proc_io
2  obj-m := ${MODNAME}.o
3  ${MODNAME}-objs := main.o
4
5  KDIR := /lib/modules/$(shell uname -r)/build
6
7  all:
8      $(MAKE) -C $(KDIR) M=$(PWD) modules
9
10 clean:
11     $(MAKE) -C $(KDIR) M=$(PWD) clean

```

Figure 15: Makefile (no warnings)

```

gvk@gvk:~/Desktop/EE516/PR02/task04 (master) $ make
make -C /lib/modules/4.4.0-38-generic/build M=/home/gvk/Desktop/EE516/PR02/task04 modules
make[1]: Entering directory '/usr/src/linux-headers-4.4.0-38-generic'
CC [M] /home/gvk/Desktop/EE516/PR02/task04/main.o
LD [M] /home/gvk/Desktop/EE516/PR02/task04/proc_io.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/gvk/Desktop/EE516/PR02/task04/proc_io.mod.o
LD [M] /home/gvk/Desktop/EE516/PR02/task04/proc_io.ko
make[1]: Leaving directory '/usr/src/linux-headers-4.4.0-38-generic'

```

```

gvk@gvk:~/Desktop/EE516/PR02/task04 (master) $ cat /proc/proc_io | grep colord
1015   colord      859003      11075      12020      1470      10805248      0      0
gvk@gvk:~/Desktop/EE516/PR02/task04 (master) $ sudo cat /proc/1015/io
rchar: 859003
wchar: 11075
syscr: 12020
syscw: 1470
read_bytes: 10805248
write_bytes: 0
cancelled_write_bytes: 0

```

Figure 16: Verification (colord)

Task 5: Sorting Features

Create a process monitoring tool with sorting feature

task05

- main.c
- Makefile
- pm_list.c
- pm_list.h
- sort.c
- sort.h
- utils.h

The task is divided into 2 modules:

1. **Sorting Module** – It is used to view or update current sorting order
2. **PM (procmon) Module** – It is used to create process monitoring

```
1 #pragma once
2
3 enum {
4     SORT_ORDER_PID = 0,
5     SORT_ORDER_VIRT,
6     SORT_ORDER_RSS,
7     SORT_ORDER_IO
8 };
9
10 /* Initializes sorting module
11  * Return:
12  *   < 0 on error
13  */
14 int
15 sort_module_init(void);
16
17 /* De-initializes sorting module
18  */
19 void
20 sort_module_exit(void);
21
22 /* Returns current sorting order
23  * Refer SORT_ORDER_<xyz> enumeration for return values
24  */
25 inline int
26 sort_get_order(void);
```

```
88 /* file operations */
89 static struct file_operations sort_ops = {
90     .owner = THIS_MODULE,
91     .open = sort_open,
92
93     /* read method for sequential files */
94     .read = seq_read,
95
96     /* llseek method for sequential files */
97     .llseek = seq_lseek,
98
99     /* write() system call on VFS */
100    .write = sort_write,
101
102    /* free the structures associated with sequential file */
103    .release = single_release,
104 };
105
106 void
107 sort_module_exit(void)
108 {
109     dbg("");
110
111     if (ps != NULL)
112         proc_remove(ps);
113 }
114
115 int
116 sort_module_init(void)
117 {
118     dbg("");
119
120     /* create /proc/procmon_sorting */
121     ps = proc_create(PROC_SORT, 0, NULL, &sort_ops);
122     if (ps == NULL) {
123         err("Failed to create procmon_sorting");
124         goto error;
125     }
126     return 0;
127
128 error:
129     sort_module_exit();
130     return -1;
131 }
132
133 inline int
134 sort_get_order(void)
135 {
136     return sort_order;
137 }
```

Figure 17: Sorting Module APIs exposed to PM module

```

48 static ssize_t
49 sort_write(struct file *file, const char __user *user_buf, size_t length, loff_t *offset)
50 {
51     char buf[BUF_SIZE];
52     dbg("");
53
54     /* clear buffer */
55     memset(buf, 0x00, sizeof(buf));
56
57     /* resize */
58     if (length > BUF_SIZE)
59         length = BUF_SIZE;
60
61     /* copy data from user space */
62     if (copy_from_user(buf, user_buf, length))
63         return -EFAULT; /* Bad address */
64
65     /* null terminate buffer */
66     buf[length - 1] = '\0';
67
68     dbg("buf: [%s]", buf);
69
70     /* set sorting order */
71     if (strcmp(buf, "pid") == 0)
72         sort_order = SORT_ORDER_PID;
73     else if (strcmp(buf, "virt") == 0)
74         sort_order = SORT_ORDER_VIRT;
75     else if (strcmp(buf, "rss") == 0)
76         sort_order = SORT_ORDER_RSS;
77     else if (strcmp(buf, "io") == 0)
78         sort_order = SORT_ORDER_IO;
79     else
80         info("Invalid value! sort_order is not changed");
81
82     dbg("sort_order: [%d]", sort_order);
83
84     /* copied from user space, return bytes */
85     return length;
86 }

```

Figure 18: Sorting Module write() file operation to update sort_order

```

15
16 static int
17 sort_show(struct seq_file *m, void *v)
18 {
19     switch (sort_order) {
20     case SORT_ORDER_VIRT:
21         seq_printf(m, "PID \t [VIRT] \t RSS \t I/O\n");
22         break;
23
24     case SORT_ORDER_RSS:
25         seq_printf(m, "PID \t VIRT \t [RSS] \t I/O\n");
26         break;
27
28     case SORT_ORDER_IO:
29         seq_printf(m, "PID \t VIRT \t RSS \t [I/O]\n");
30         break;
31
32     case SORT_ORDER_PID:
33     default: /* fallthrough */
34         seq_printf(m, "[PID] \t VIRT \t RSS \t I/O\n");
35         break;
36     }
37
38     return 0;
39 }
40
41 static int
42 sort_open(struct inode *inode, struct file *file)
43 {
44     dbg("");
45     return single_open(file, sort_show, NULL);
46 }

```

Figure 19: Sorting Module open() file operation

```

52  /* file operations */
53  static struct file_operations fops_pm = {
54      .owner = THIS_MODULE,
55      .open = pm_open,
56
57      /* read method for sequential files */
58      .read = seq_read,
59
60      /* llseek method for sequential files */
61      .llseek = seq_lseek,
62
63      /* free the structures associated with sequential file */
64      .release = pm_release,
65  };
66
67  static void
68  _pm_module_exit(void)
69  {
70      dbg("");
71
72      /* de-initialize sorting module */
73      sort_module_exit();
74
75      /* remove procmon */
76      if (pm != NULL)
77          proc_remove(pm);
78  }
79
80  static int __init
81  pm_module_init(void)
82  {
83      int ret;
84      dbg("");
85
86      /* create /proc/procmon */
87      pm = proc_create(PROC_NAME, 0, NULL, &fops_pm);
88      if (pm == NULL) {
89          err("Failed to create procmon");
90          goto error;
91      }
92
93      /* initialize sorting module */
94      ret = sort_module_init();
95      if (ret < 0) {
96          err("Failed to initialize sorting module");
97          goto error;
98      }
99      return 0;
100
101  error:
102      _pm_module_exit();
103      return -1;
104  }
105
106  static void __exit
107  pm_module_exit(void)
108  {
109      _pm_module_exit();
110  }

```

Figure 20: PM Module init / exit
(Sorting module is init/exit by PM module)

```

12 static int
13 pm_show(struct seq_file *m, void *v)
14 {
15     dbg("");
16
17     /* show pm_list */
18     return pm_list_show(m);
19 }
20
21 static int
22 pm_open(struct inode *inode, struct file *file)
23 {
24     int ret;
25     dbg("");
26
27     /* initialize pm_list
28      * it means to parse all processes & save them in
29      * linked list owned by this module
30      */
31     ret = pm_list_init();
32     if (ret < 0) {
33         err("Failed to initialize pm_list: %d", ret);
34         return ret;
35     }
36
37     return single_open(file, pm_show, NULL);
38 }
39
40 static int
41 pm_release(struct inode *inode, struct file *file)
42 {
43     dbg("");
44
45     /* clean-up pm_list */
46     pm_list_deinit();
47
48     /* release sequence file */
49     return single_release(inode, file);
50 }
51

```

Figure 21: PM Module open() and release()
(**pm_list** is PM Module's data structure discussed in next page)

```

7  /* pm_list node entry */
8  struct pm_list_entry {
9      struct list_head entries;
10
11     pid_t pid;
12     char name[TASK_COMM_LEN];
13     unsigned long long virt; /* in KB */
14     long long rss; /* in KB */
15     unsigned long long disk_read; /* in KB */
16     unsigned long long disk_write; /* in KB */
17 };
18
19 /* first entry of pm_list */
20 extern struct pm_list_entry pm_list_init_entry;
21
22 /* macro to find next pm_list element */
23 #define pm_list_next_entry(e) \
24     container_of((e)->entries.next, struct pm_list_entry, entries)
25
26 /* initializes pm_list
27  * Return:
28  *   < 0 on error
29  */
30 int pm_list_init(void);
31
32 /* deinit pm_list
33  */
34 void pm_list_deinit(void);
35
36 /* shows pm_list after sorting in
37  * currently set order
38  */
39 int pm_list_show(struct seq_file *m);

```

Figure 22: pm_list data structure & APIs.
(Each API is discussed next)

Notes:

1. The design of this data structure closely resembles task_struct from the Linux Kernel.

2. Reference:

https://github.com/torvalds/linux/blob/master/init/init_task.c

3. “struct list_head” is the kernel way of providing linked lists, which has been referred from

<https://github.com/torvalds/linux/blob/master/include/linux/list.h>


```

125 int
126 pm_list_init(void)
127 {
128     struct task_struct *tsk, *t;
129     struct pm_list_entry *entry;
130     char name[TASK_COMM_LEN];
131     unsigned long long virt;
132     long long rss;
133     struct task_io_accounting acct;
134
135     dbg("");
136
137     /* Add to linked list */
138     for_each_process(tsk) {
139         dbg("[ADD] name: [%s] pid: [%d]", get_task_comm(name, tsk),
140             task_pid_nr(tsk));
141
142         /* Calculate memory */
143         virt = 0;
144         rss = 0;
145         if (tsk->active_mm != NULL) {
146             virt = tsk->active_mm->total_vm;
147             virt *= (PAGE_SIZE >> 10); /* convert pages to KB */
148
149             rss += atomic_long_read(&tsk->active_mm->rss_stat.count[MM_FILEPAGES]);
150             rss += atomic_long_read(&tsk->active_mm->rss_stat.count[MM_ANONPAGES]);
151             rss *= (PAGE_SIZE >> 10); /* convert pages to KB */
152         }
153
154         /* I/O */
155         acct = tsk->ioac;
156         t = tsk;
157
158         /* Account for each thread */
159         task_io_accounting_add(&acct, &tsk->signal->ioac);
160         while_each_thread(tsk, t)
161             task_io_accounting_add(&acct, &t->ioac);
162
163         /* Allocate Linked List node */
164         entry = alloc_pm_list_node(task_pid_nr(tsk),
165             get_task_comm(name, tsk),
166             virt, rss,
167             (acct.read_bytes >> 10), /* divide by 1024 to convert into KB */
168             (acct.write_bytes >> 10));
169         if (entry == NULL) {
170             info("Skipping PID: [%d]", task_pid_nr(tsk));
171             continue;
172         }
173
174         /* Add to List */
175         list_add(&entry->entries, &pm_list_init_entry.entries);
176     }
177
178     return 0;
179 }

```

Figure 23: pm_list_init() – API for creating & initializing a PM list

```

56 int
57 pm_list_show(struct seq_file *m)
58 {
59     struct list_head *cursor, *temp;
60     struct pm_list_entry *entry;
61     dbg("");
62
63     /* Sort list */
64     list_sort(NULL,
65              &pm_list_init_entry.entries,
66              sort_pm_list_entries);
67
68     /* title */
69     seq_printf(m, "===== Process Monitoring Manager for EE516 =====\n");
70
71     /* header */
72     seq_printf(m, "PID          ProcessName          VIRT(KB)          "
73              "RSS Mem(KB)          DiskRead(KB)          DiskWrite(KB)          "
74              "Total I/O(KB)          \n");
75
76     list_for_each_safe(cursor, temp, &pm_list_init_entry.entries) {
77         /* typecast */
78         entry = list_entry(cursor, struct pm_list_entry, entries);
79
80         /* print information */
81         seq_printf(m, "%-10u%-20s%-20llu%-20llu%-20llu%-20llu%-20llu\n",
82                  entry->pid,
83                  entry->name,
84                  entry->virt,
85                  entry->rss,
86                  entry->disk_read,
87                  entry->disk_write,
88                  (entry->disk_read + entry->disk_write));
89     }
90
91     return 0;
92 }

```

```

23 /* sort function for pm_list */
24 static int
25 sort_pm_list_entries(void *priv, struct list_head *a, struct list_head *b)
26 {
27     struct pm_list_entry *entry_a, *entry_b;
28     int sort_order;
29
30     /* typecast */
31     entry_a = list_entry(a, struct pm_list_entry, entries);
32     entry_b = list_entry(b, struct pm_list_entry, entries);
33
34     /* find current sort order */
35     sort_order = sort_get_order();
36
37     /* sort */
38     switch (sort_order) {
39     case SORT_ORDER_VIRT:
40         return (entry_b->virt - entry_a->virt);
41
42     case SORT_ORDER_RSS:
43         return (entry_b->rss - entry_a->rss);
44
45     case SORT_ORDER_IO:
46         /* subtract & then add to prevent overflow */
47         return ((entry_b->disk_write - entry_a->disk_write) + \
48                 (entry_b->disk_read - entry_a->disk_read));
49
50     case SORT_ORDER_PID:
51     default:
52         return (entry_a->pid - entry_b->pid);
53     }
54 }

```

Figure 24: pm_list_show() – API for displaying PM list in sorted order

```

181 void
182 pm_list_deinit(void)
183 {
184     struct list_head *cursor, *temp;
185     struct pm_list_entry *entry;
186     dbg("");
187
188     1 list_for_each_safe(cursor, temp, &pm_list_init_entry.entries) {
189         entry = list_entry(cursor, struct pm_list_entry, entries);
190         dbg("[REMOVE] name: [%s] pid: [%d]", entry->name, entry->pid);
191
192         2 list_del(cursor);
193         free_pm_list_node(entry);
194     }
195
196     dbg("Is empty? %d", list_empty(&pm_list_init_entry.entries));
197 }

```

Figure 25: pm_list_deinit() – API for freeing a PM list

```

1  MODNAME := procmon
2  obj-m := ${MODNAME}.o
3  ${MODNAME}-objs := main.o pm_list.o sort.o
4
5  KDIR := /lib/modules/$(shell uname -r)/build
6
7  all:
8      $(MAKE) -C $(KDIR) M=$(PWD) modules
9
10 clean:
11     $(MAKE) -C $(KDIR) M=$(PWD) clean

```

Figure 26: Makefile (no warnings)

```

gvk@lra@gvkalra-desktop ~/Desktop/EE516/PR02/task05 (master) $ make
make -C /lib/modules/4.4.0-38-generic/build M=/home/gvkalra/Desktop/EE516/PR02/task05 modules
make[1]: Entering directory '/usr/src/linux-headers-4.4.0-38-generic'
CC [M] /home/gvkalra/Desktop/EE516/PR02/task05/main.o
CC [M] /home/gvkalra/Desktop/EE516/PR02/task05/pm_list.o
CC [M] /home/gvkalra/Desktop/EE516/PR02/task05/sort.o
LD [M] /home/gvkalra/Desktop/EE516/PR02/task05/procmon.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/gvkalra/Desktop/EE516/PR02/task05/procmon.mod.o
LD [M] /home/gvkalra/Desktop/EE516/PR02/task05/procmon.ko
make[1]: Leaving directory '/usr/src/linux-headers-4.4.0-38-generic'

```

```

root@gvkalra-desktop:~# cat /proc/procmon_sorting
[PID] VIRT RSS I/O
root@gvkalra-desktop:~# echo virt > /proc/procmon_sorting
root@gvkalra-desktop:~# cat /proc/procmon
===== Process Monitoring Manager for EE516 =====

```

PID	ProcessName	VIRT(KB)	RSS	Mem(KB)	DiskRead(KB)	DiskWrite(KB)	Total I/O(KB)
3454	synergys	3214860	21156		1252	0	1252
2372	evolution-calen	2196228	75640		600	256	856
2324	compiz	1849444	525492		24892	128	25020
2529	slack	1685444	146980		49980	12604	62584
21950	chrome	1496680	190064		0	0	0
2736	slack	1496056	345404		13428	84	13512

Figure 27: Verification (VIRT)

```

root@gvkalra-desktop:~# echo rss > /proc/procmon_sorting
root@gvkalra-desktop:~# cat /proc/procmon
===== Process Monitoring Manager for EE516 =====

```

PID	ProcessName	VIRT(KB)	RSS	Mem(KB)	DiskRead(KB)	DiskWrite(KB)	Total I/O(KB)
2324	compiz	1849444	525492		24892	128	25020
2736	slack	1496056	345404		13428	84	13512
2780	slack	1379216	264996		4	20	24
21721	chrome	1188368	218420		1180	603128	604308
22011	chrome	980040	192536		0	0	0
21950	chrome	1496680	190180		0	0	0

Figure 28: Verification (RSS)

Problem 1

Describe similarity and difference Process (Thread Group) and Thread in the Linux

In Linux, Processes & Threads are all represented by `task_struct`. The `pid` field of `task_struct` is unique for every process & thread.

However, since POSIX mandates that all threads of a process should have the same ProcessID (as seen by the user), Linux uses “Thread Group ID” (TGID) to satisfy POSIX mandate.

A TGID is the PID of the thread that started the whole process.

```
1474
1475 struct task_struct {
1591
1592     pid_t pid;
1593     pid_t tgid;
1594
```

Figure 29: pid/tgid in `task_struct`.

Source: <https://github.com/torvalds/linux/blob/master/include/linux/sched.h>

In Linux, both processes & threads are created using `clone()` system call. The ‘flags’ passed to `clone()` specify the degree of data sharing. Threads in Linux have it’s own stack but shares Heap, BSS, Data and Text with other threads having same TGID.

The CFS (Completely Fair Scheduler) as used in Kernel doesn’t discriminate between a thread & a process. It works on “tasks” and is not aware of thread & process abstractions.

Problem 2

Use of `copy_to_user` / `copy_from_user` functions in `proc` filesystem

In Linux, user-space is separated from kernel-space in all aspects. To arbitrate data between user / kernel space, Linux provides two functions:

1. **`copy_to_user`**: Copy data from kernel space to user space
2. **`copy_from_user`**: Copy data from user space to kernel space

A `proc` file system (`procfs`) is just another type of file system on top of Virtual File System (VFS) & thus applies the same rule of data exchange between user & kernel space.